



ĐẠI HỌC ĐÀ NẴNG

TRƯỜNG ĐẠI HỌC CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG VIỆT - HÀN
VIETNAM - KOREA UNIVERSITY OF INFORMATION AND COMMUNICATION TECHNOLOGY

한-베정보통신기술대학교

Data Warehouse

Faculty of Computer Science



ĐẠI HỌC ĐÀ NẴNG

TRƯỜNG ĐẠI HỌC CÔNG NGHỆ THÔNG TIN VÀ TRUYỀN THÔNG VIỆT - HÀN
VIETNAM - KOREA UNIVERSITY OF INFORMATION AND COMMUNICATION TECHNOLOGY

한-베정보통신기술대학교

Chapter 6: Extraction, Transformation, and Loading (ETL)

Chapter Content

- ◆ Business Process Modeling Notation
- ◆ Conceptual ETL Design using BPMN
- ◆ Conceptual Design of the Northwind ETL process
- ◆ Integration Services and Kettle
- ◆ The Northwind ETL process in Integration Services
- ◆ The Northwind ETL process in Kettle

Extraction, Transformation, and Loading (ETL)

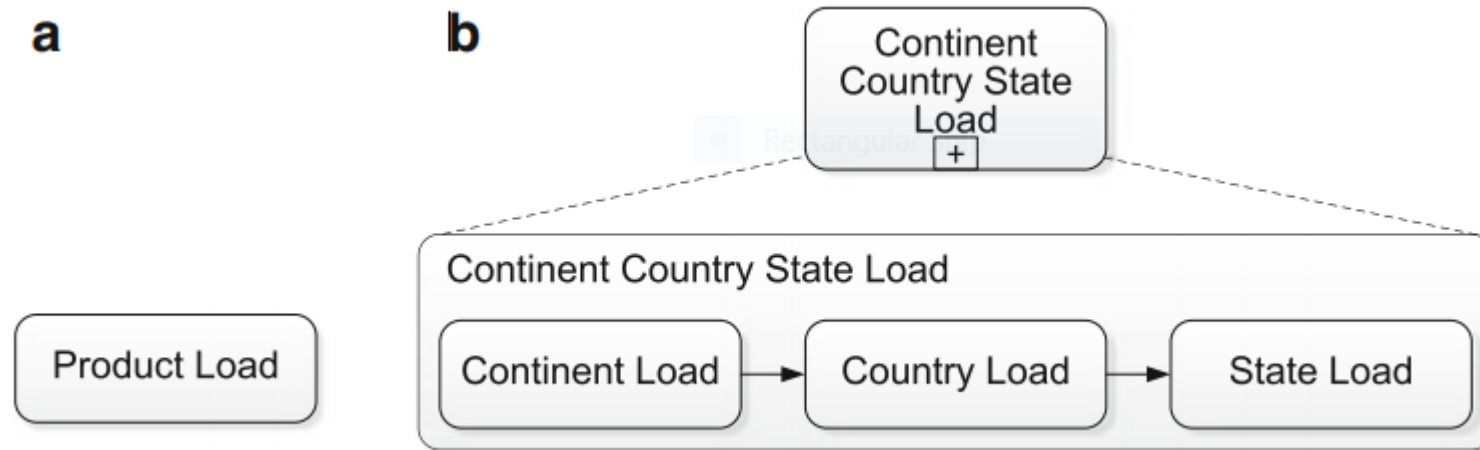
- ◆ **Extract** data from internal and external sources, **transform** data, and **load** data into a data warehouse
- ◆ No agreed way to specify ETL at a conceptual level
- ◆ We study conceptual ETL design
- ◆ Conceptual model based on the Business Process Modeling Notation (BPMN)
 - Users already familiar with BPMN do not need to learn another language to design ETL
 - BPMN provides a conceptual and implementation-independent specification of processes
 - Processes expressed in BPMN can be translated into executable specifications(e.g., Microsoft's Integration Services)

Business Process Modeling Notation (BPMN)

- ◆ **Business process**: A collection of related activities or tasks whose goal is to produce a specific service or product
- ◆ **Business process modeling**: Activity of representing the business processes of an organization, so that the current processes may be analyzed and improved
- ◆ Many techniques to model business process proposed over the years
- ◆ No formal semantics for these techniques
- ◆ Formal techniques (e.g., Petri Nets): Well-defined semantics but hard to understand by business users
- ◆ A standardization process resulted in the **Business Process Modeling Notation (BPMN)** released by the Object Management Group (OMG). Current version is BPMN 2.0
- ◆ BPMN: Graphical notation for defining, understanding, and communicating the business procedures of an organization them in a standard manner
- ◆ Four basic categories of elements: **flow objects**, **connecting objects**, **swimlanes**, and **artifacts**

Flow Objects: Activities

- ◆ An **activity** is a work performed during a process
 - Can be atomic or nonatomic
 - Can be a task or a subprocess
- ◆ **Subprocess**: An encapsulated process whose details we want to hide




Activities. (a) Single task. (b) Collapsed and expanded subprocess


Flow Objects: Gateways

- ◆ Control the activity sequence in a process, based on conditions
- ◆ Represent only logic, not activities
- ◆ **Exclusive gateways** model OR-split decisions
- ◆ **Inclusive gateways** select or merge one or more flows
- ◆ **Parallel gateways** allow the synchronization between outgoing and incoming flows
 - **Splitting** parallel gateway: Analogous to an AND-split
 - **Merging** parallel gateway: Synchronizes the flow and merges all the incoming flows into a single outgoing one
- ◆ **Complex gateways** can represent complex conditions

Different types of gateways



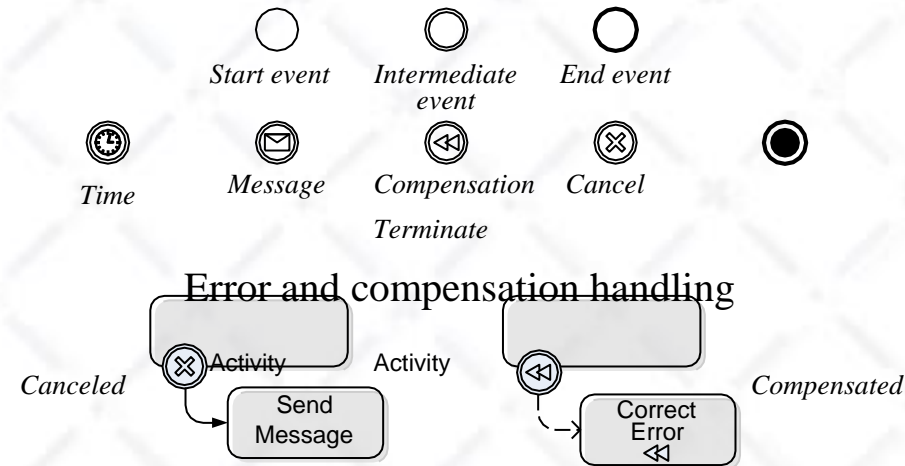
Splitting and merging gateways



Flow Objects: Events


- ◆ Represent something that happens that *affects* the sequence and timing of the workflow activities
- ◆ **Start** and **end** events indicate the beginning and ending of a process
- ◆ **Time** event: represents situations when a task must wait for some period of time before continuing
- ◆ **Message** event represents communication
- ◆ **Compensation** event represents error detection and recovery by launching compensation activities
- ◆ **Cancel** event listens to the process errors and notifies them by an explicit or implicit action
- ◆ **Terminate** When it is reached, the entire process is stopped, including all parallel processes.

Examples of events



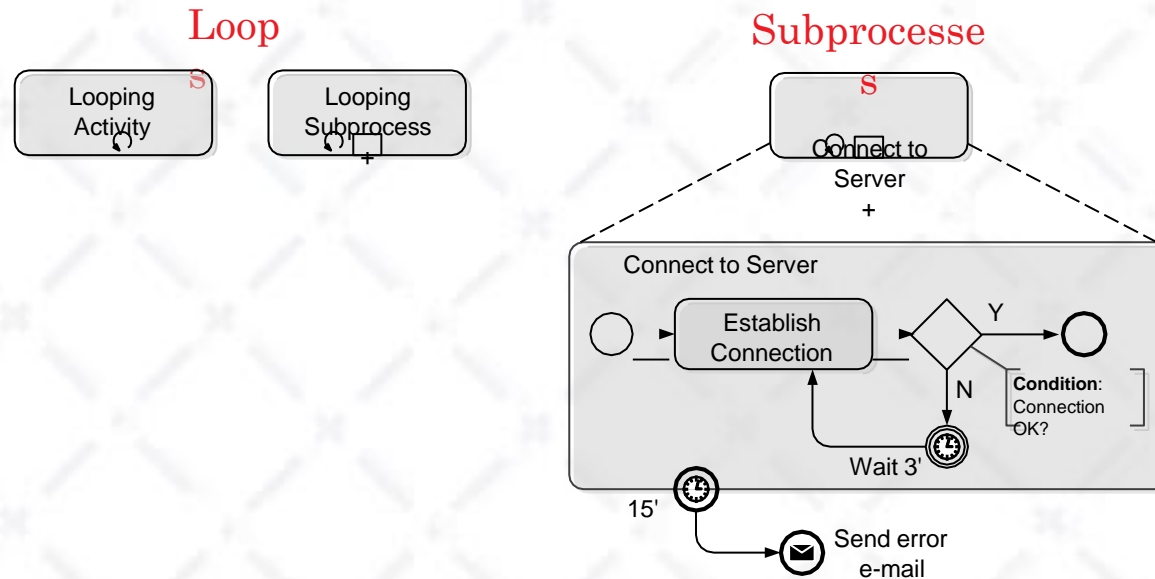
Connecting Objects

- ◆ Represent how objects are connected
- ◆ **Sequence flow**: A sequencing constraint between flow objects
 - If two activities are linked by a sequence flow, the target one starts when the source one has finished
 - If multiple sequence flows outgo from a flow object, all of them will be activated after its execution
- ◆ **Conditional sequence flow**: Adds a condition to the sequence flow
- ◆ A sequence flow may be set as the **default flow** in case of many outgoing flows (e.g., if no other condition is true in a gateway, the default flow is followed)
- ◆ A **message flow** represents the **only way** of sending and receiving of messages between pools. An **association** relates artifacts (e.g., annotations) to flow objects.



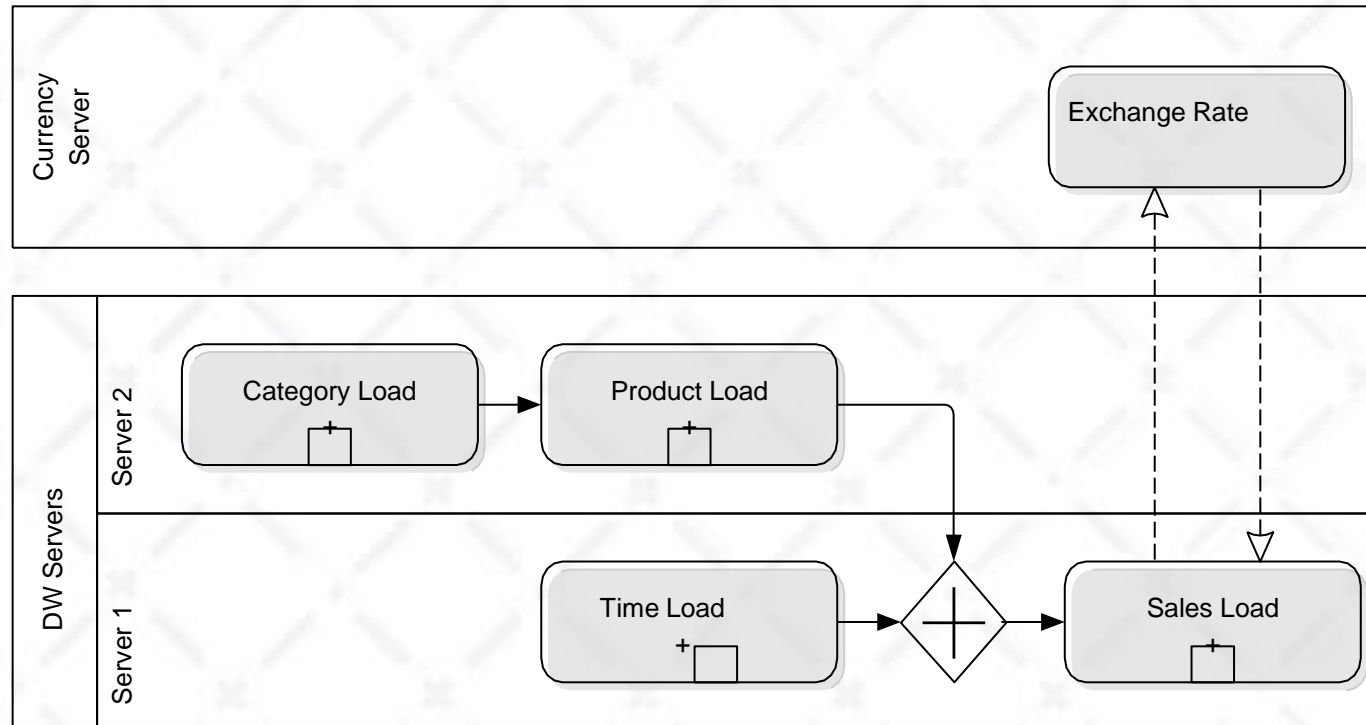
Loops and Subprocesses

- ◆ **Loop**: Execution control feature representing repeated execution of a process
- ◆ Conditions checked before or after activity. Loop ended if its condition evaluates to false.
- ◆ Figure: ETL process representing the connection to a server task
- ◆ At a high abstraction level, the **subprocess activity** hides the details
- ◆ Expansion shows details: server waits 3 minutes (**time** event). If connection not established, request launched again. If no connection after 15 minutes, task stopped, and error email sent (**message** event).



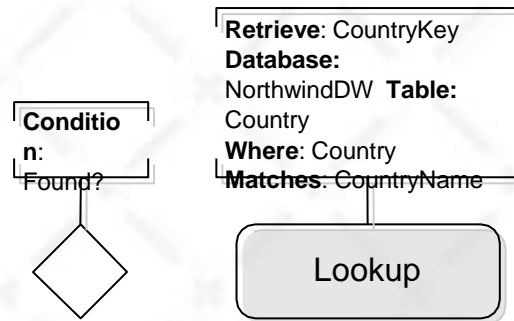
Swimlanes

- ◆ A structuring object that comprises **pools** and **lanes**
- ◆ Both allow the definition of process boundaries
- ◆ Only messages allowed between two pools, not sequence flows
- ◆ A workflow must be contained in only one pool
- ◆ One pool may be subdivided into many lanes, which represent roles or services



Artifacts

- ◆ Allow to visually represent objects outside the actual process
- ◆ Can represent data or notes that describe the process, or they can be used to organize tasks or processes
- ◆ Can be data objects, groups, and annotations
- ◆ A **data object** represents either data that are input to a process, data resulting from a process, data that needs to be collected, or data that needs to be stored.
- ◆ A **group** organizes tasks or processes that have some kind of significance in the overall model
- ◆ **Annotations** are used to express semantics about the flow objects (e.g., to indicate the attributes involved in a lookup task, or a gateway condition)



Chapter 6: Extraction, Transformation, and Loading

Outline

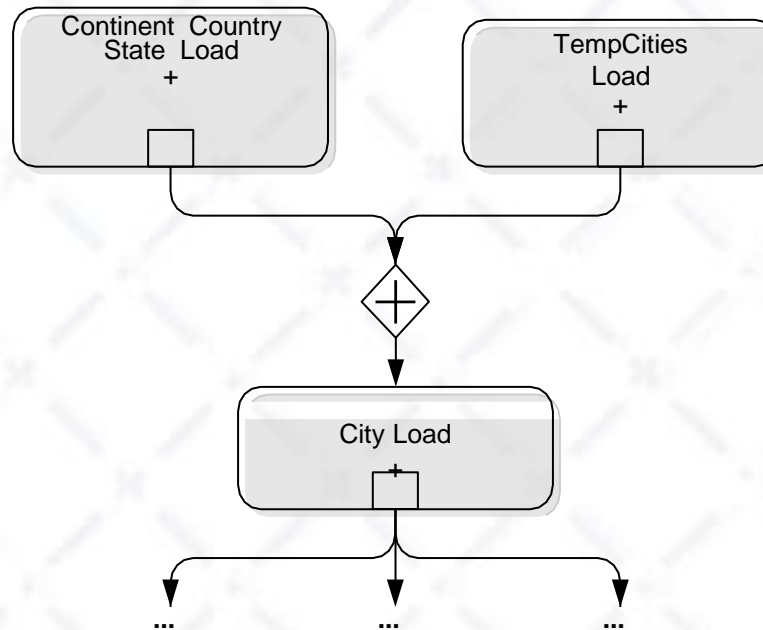
- ◆ Business Process Modeling Notation
- ② **Conceptual ETL Design using BPMN**
- ◆ Conceptual Design of the Northwind ETL
- ◆ Integration Services and Kettle
- ◆ The Northwind ETL in Integration Services
- ◆ The Northwind ETL Process in Kettle

Conceptual ETL Design using BPMN

- ◆ Basic assumption for using BPMN as conceptual model: ETL process is a type of business process
- ◆ There is no standard model for defining ETL processes
- ◆ Each tool provides its own model, too detailed to be conceptual
- ◆ Using BPMN constructs we define the most common ETL tasks and define a BPMN notation for ETL
- ◆ ETL process: A combination of **control** and **data processes**
 - Control processes manage the coarse-grained groups of tasks
 - Data processes detail how input data are transformed and output data are produced
- ◆ Two kinds of tasks in ETL conceptual modeling
 - **Control tasks** highlight the control procedures provided by BPMN. Represent a **workflow** (arrows represent the precedence between activities)
 - **Data tasks** refer to the tasks that directly manipulate data during an ETL process. Represent a **data flow** (arrows represent data 'flowing' along them)

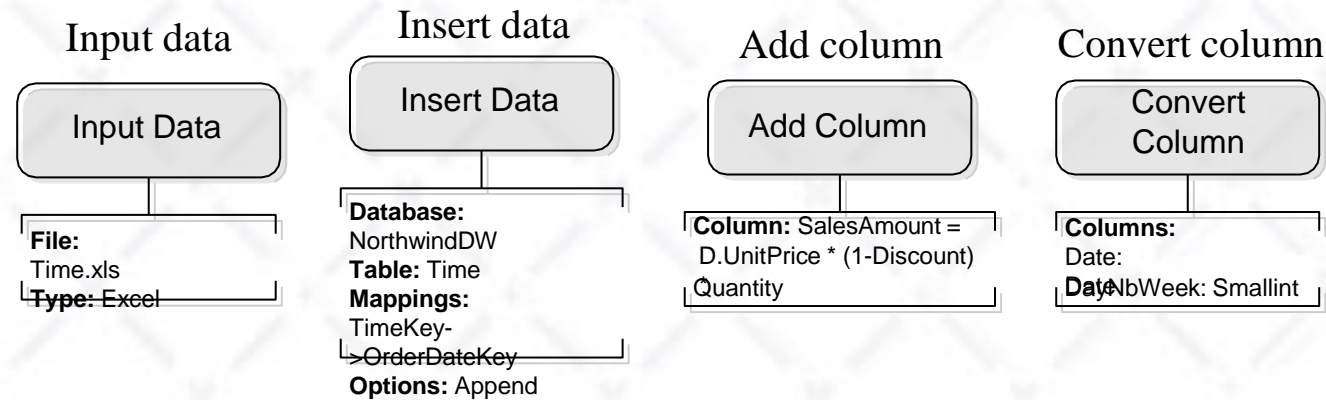
Control Tasks

- ◆ Represent the workflow sequence or **orchestration** of the ETL process independently of the data flow
- ◆ Control tasks are represented by means of BPMN constructs described
- ◆ For example, gateways are used to control the sequence of activities in an ETL process
- ◆ The most used types of gateways in an ETL context are exclusive and parallel

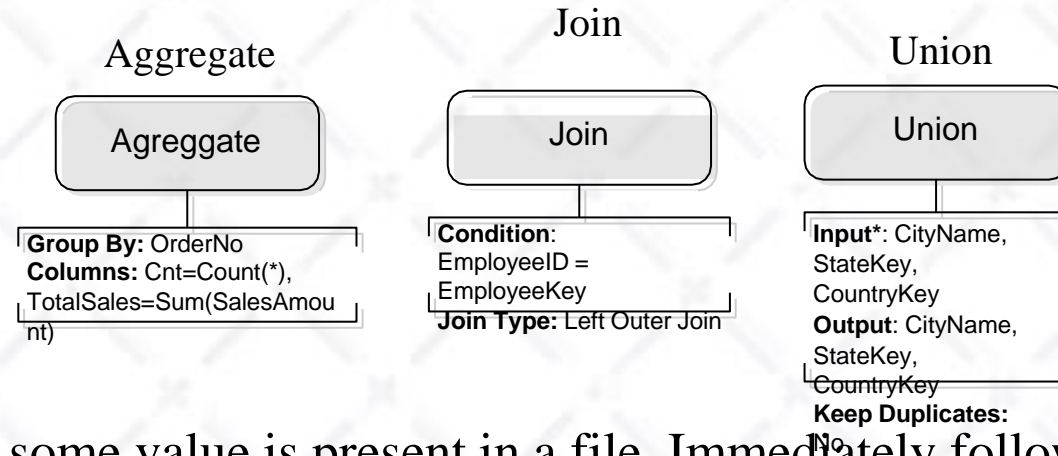


Data Tasks

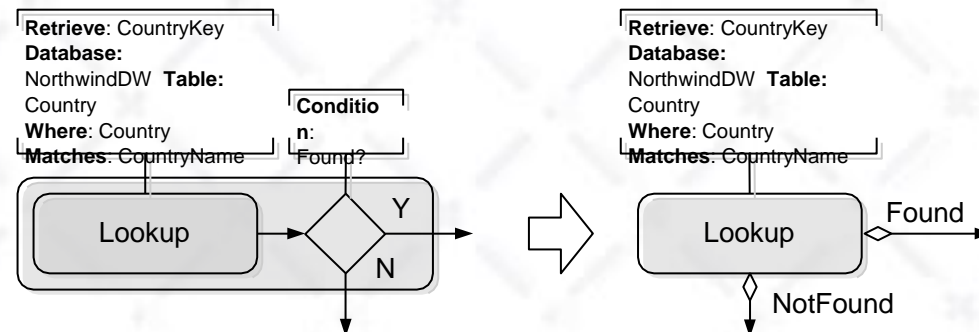
- ◆ Show how data are manipulated **within** an activity
- ◆ At lower abstraction level than control tasks
- ◆ Represent activities typically carried out to manipulate data: input and output data, data conversion and transformation (for instance, change the data type of an attribute, add a column, remove duplicates, and so on)
- ◆ We denote these tasks **unary data tasks** since they receive one input flow
- ◆ ***n*-ary data tasks** receive as input more than one flow (e.g., this is the case of union, join, *difference*,...)
- ◆ **Row operations** are transformations applied to the source or target data on a row-by-row basis, e.g., updating the value of a column
- ◆ **Rowset operations** deal with a set of rows, e.g., aggregation is a rowset operation



Rowset Data Tasks



Lookup Data Tasks check if some value is present in a file. Immediately followed by an exclusive gateway with a branching condition. We use a shorthand replacing these two tasks by 2 conditional flows.



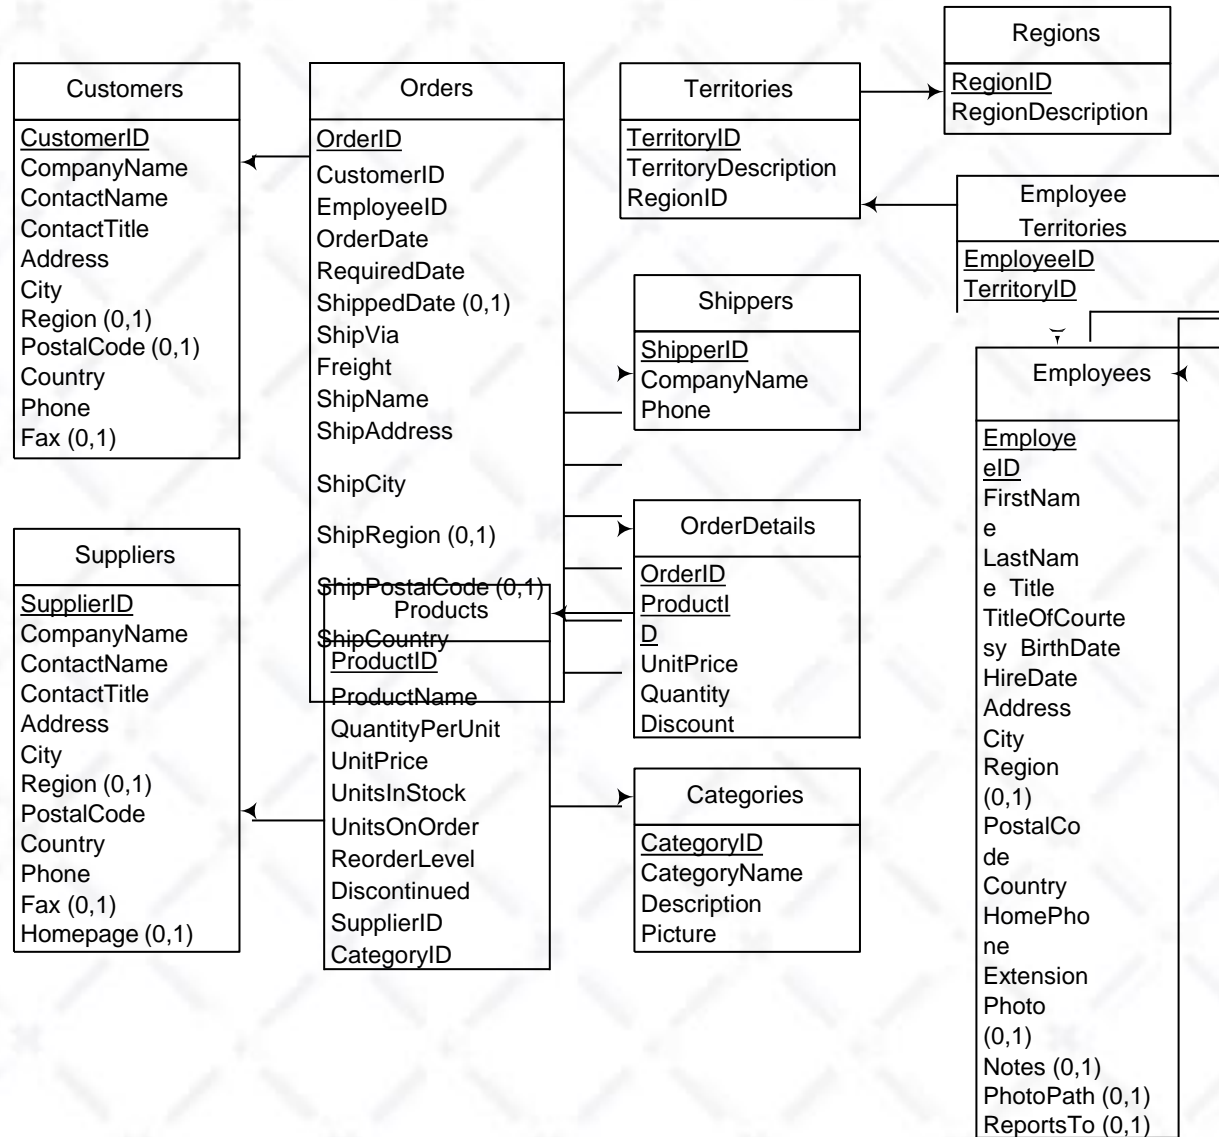
Shorthand notation for the lookup task

Chapter 6: Extraction, Transformation, and Loading

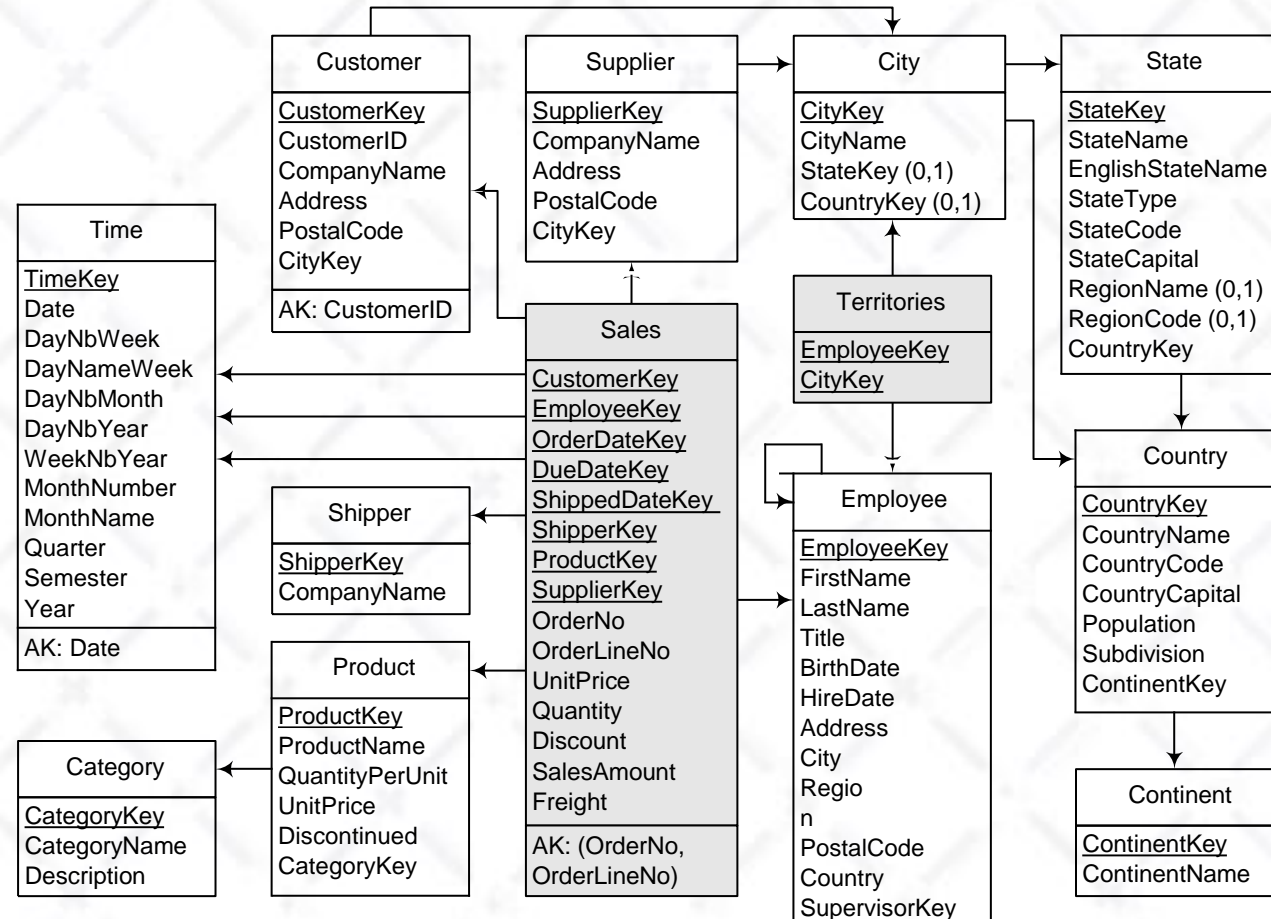
Outline

- ◆ Business Process Modeling Notation
- ◆ Conceptual ETL Design using BPMN
- ② **Conceptual Design of the Northwind ETL process**
- ◆ Integration Services and Kettle
- ◆ The Northwind ETL process in Integration Services
- ◆ The Northwind ETL process in Kettle

Schema of the Northwind Operational Database



Schema of the Northwind Data Warehouse



Conceptual Design of the Northwind ETL: Data Sources

```
<?xml version="1.0" encoding="ISO-8859-1"?>
```

```
<Continents>
```

```
<Continent>
```

```
<ContinentName>Europe</ContinentName>
```

```
<Country>
```

```
<CountryName>Austria</CountryName>
```

```
<CountryCode>AT</CountryCode>
```

```
<CountryCapital>Vienna</CountryCapital>
```

```
<Population>8316487</Population>
```

```
<Subdivision>Austria is divided into nine Bundeslnder, or simply Lnder  
(states; sing. Land).</Subdivision>
```

```
<State type="state">
```

```
<StateName>Burgenland</StateName>
```

```
<StateCode>BU</StateCode>
```

```
<StateCapital>Eisenstadt</StateCapital>
```

```
</State>
```

```
<State type="state">
```

```
<StateName>Krnnten</StateName>
```

```
<StateCode>KA</StateCode>
```

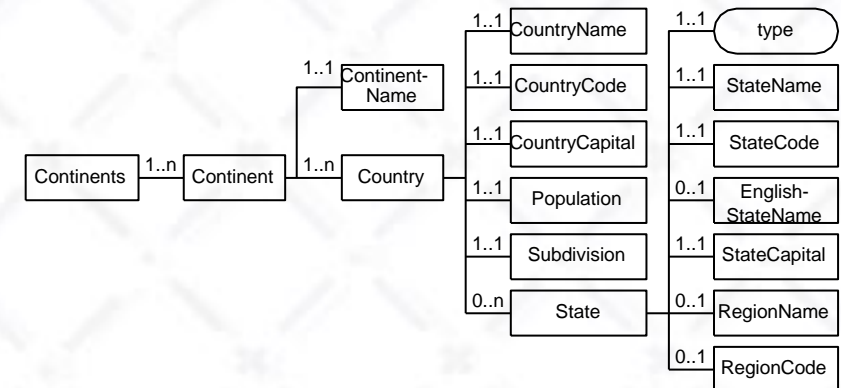
```
<EnglishStateName>Carinthia</EnglishStateName>
```

```
<StateCapital>Klagenfurt</StateCapital>
```

```
</State>
```

```
...
```

- ◆ File **Time.xls** contains data for loading the **Time** dimension, spanning the dates in table **Orders** of the operational database
- ◆ Dimensions **Customer** and **Supplier** share the geographic hierarchy starting at the **City** level
- ◆ Data for the hierarchy **State** → **Country** → **Continent** loaded from **Territories.xml**



XML Schema of **Territories.xml**

Start of the file **Territories.xml**

Conceptual Design of the Northwind ETL: Data Sources

- ◆ File called **Cities.txt** identifies to which state or province a city belongs
- ◆ Contains three fields separated by tabs and begins as shown below
- ◆ For cities located in countries that do not have states (e.g., Singapore), second field is set to null
- ◆ The file is also used to identify to which state corresponds the city in the attribute **TerritoryDescription** of table **Territories**

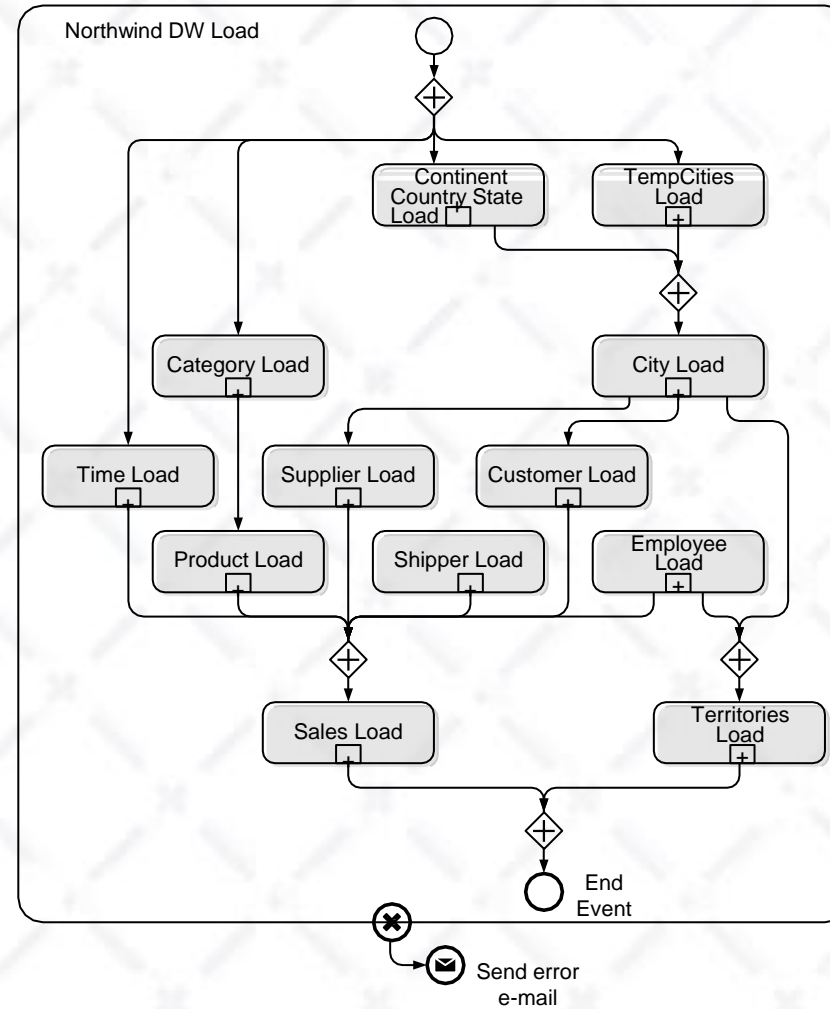
City → State → Country
Aachen → North Rhine-Westphalia → Germany
Albuquerque → New Mexico → USA
Anchorage → Alaska → USA
Ann Arbor → Michigan → USA
Annecy → Haute-Savoie → France

... Beginning of the file **Cities.txt**

| TempCities |
|------------|
| City |
| State |
| Country |


Associated table **TempCities**

Conceptual Design of the Northwind ETL: Overall View



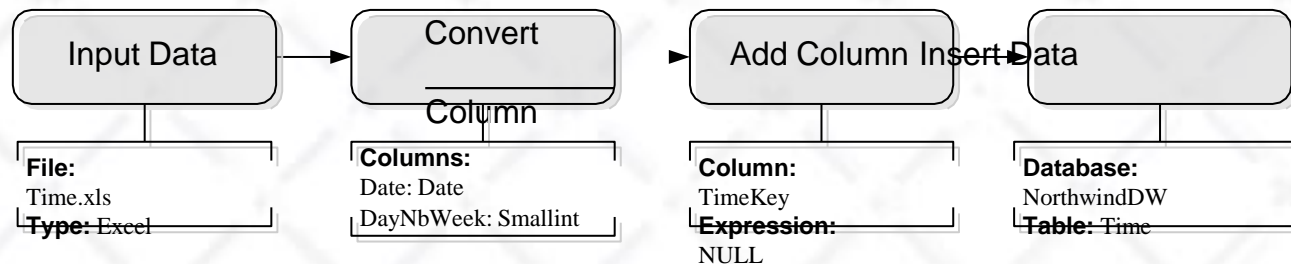
Conceptual Design of the Northwind ETL

◆ Load of the **Category** dimension table



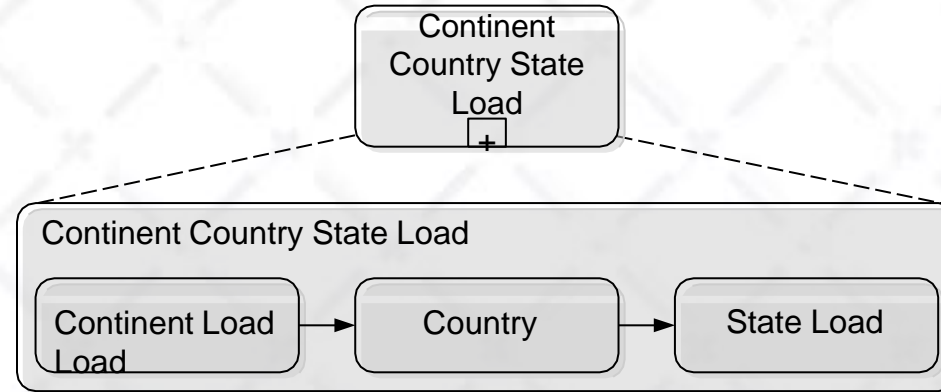
- Input task loads table **Categories** from the operational database
- Insert task loads the table **Category** in the data warehouse, mapping **CategoryID** to **CategoryKey** attribute in the **Category** table

◆ Loading the **Time** dimension table from an Excel file is similar, but includes a data type conversion, and an addition of the column **TimeKey**




Conceptual Design of the Northwind ETL

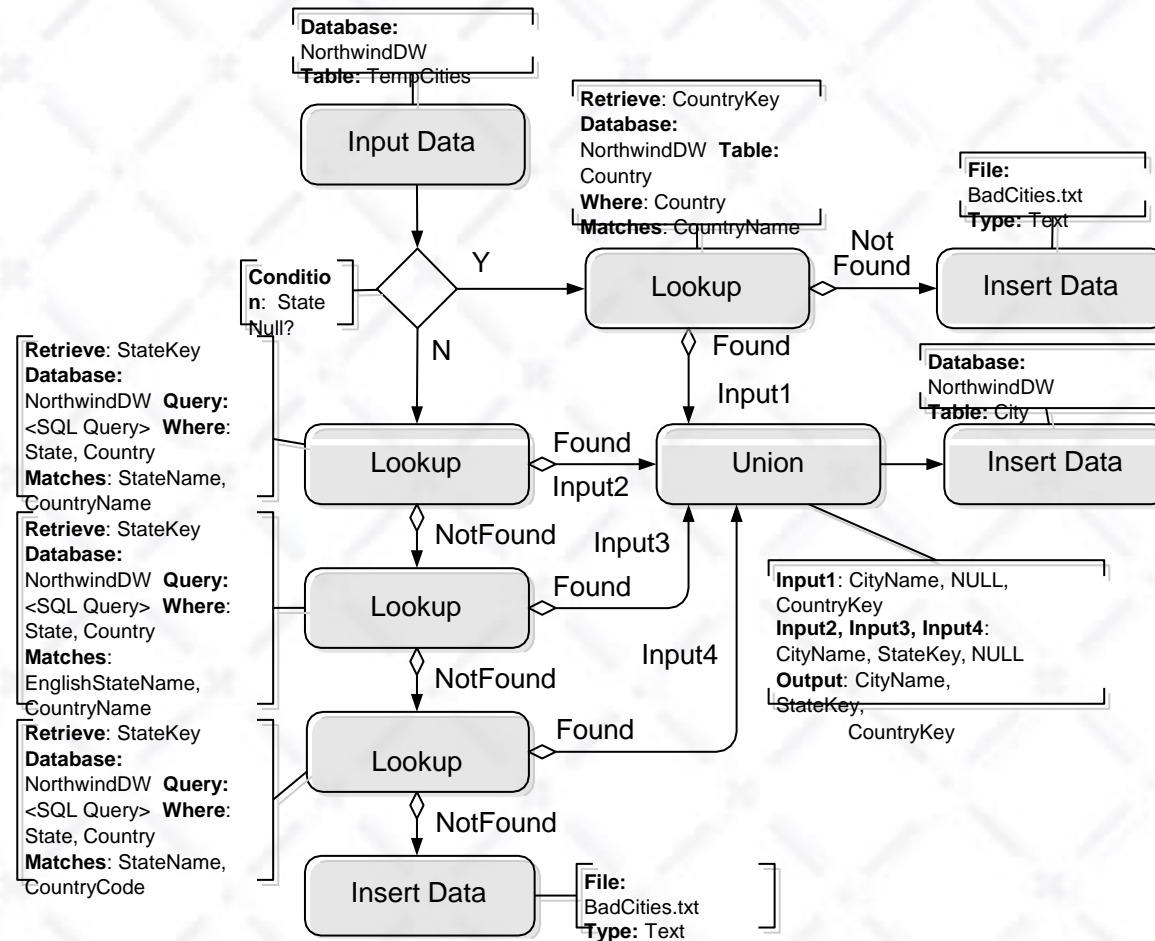
- ◆ Loading the **City** level first requires loading the **Geography** hierarchy **State** → **Country** → **Continent**
- ◆ Associated control task



- ◆ Load of the **Continent** table



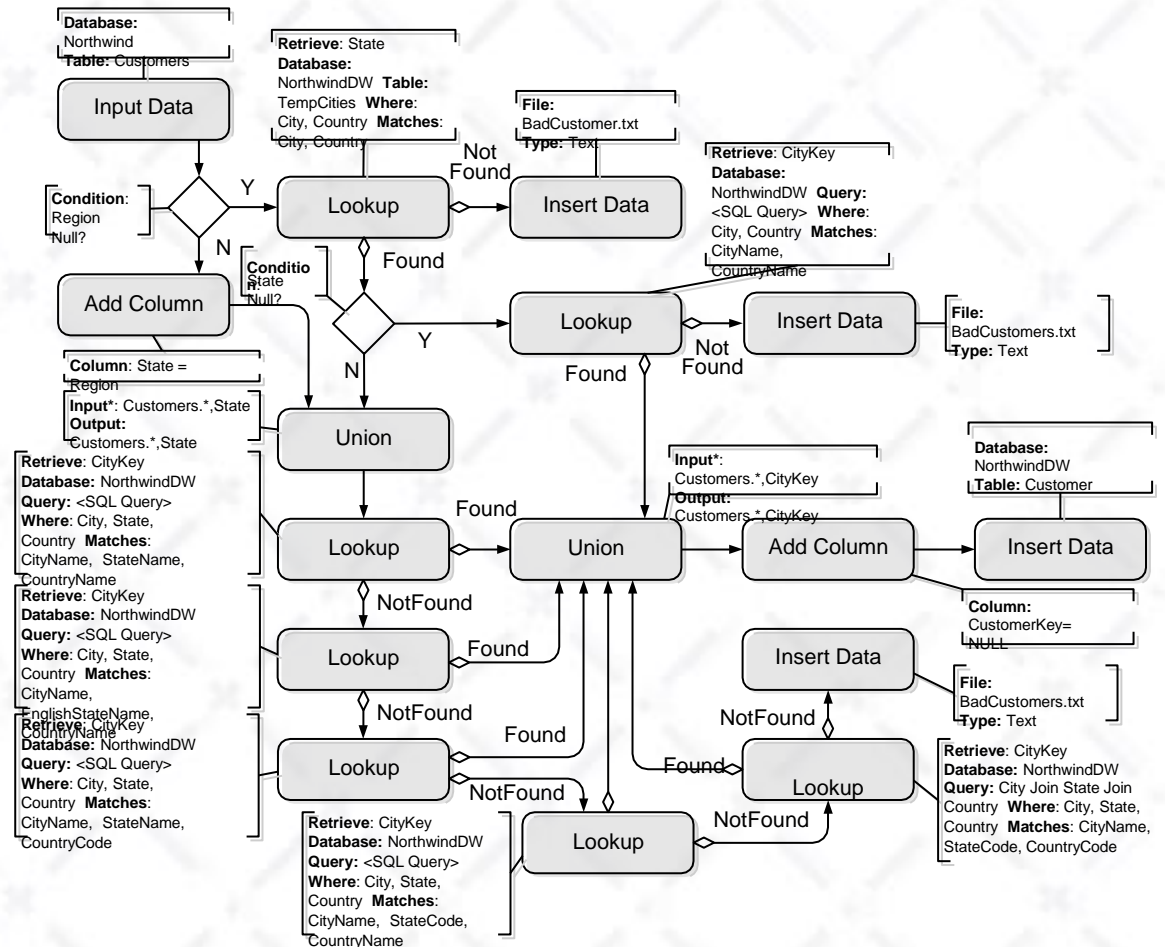
Load of the City Level



Load of the City Level

- ◆ Assume that a table **TempCities(City,State,Country)** has been created and populated from **Cities.txt**
- ◆ First task is an input data over **TempCities**
- ◆ An exclusive gateway tests whether **State** is null or not
 - If so, lookup obtains the **CountryKey**
 - If not, we match (**State, Country**) pairs in **TempCities** to values in the **State** and **Country** tables
- ◆ Finally, union performed with the results of the four flows, and table is loaded with an insert data task
- ◆ Records for which the state and/or country are not found are stored into a **BadCities.txt** file.

Load of the Customer Level

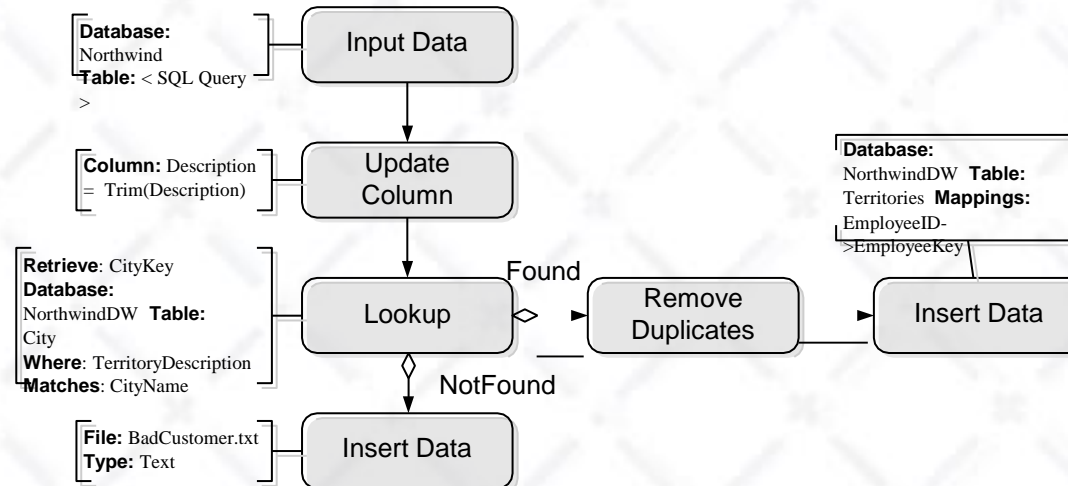


Load of the Customer Level

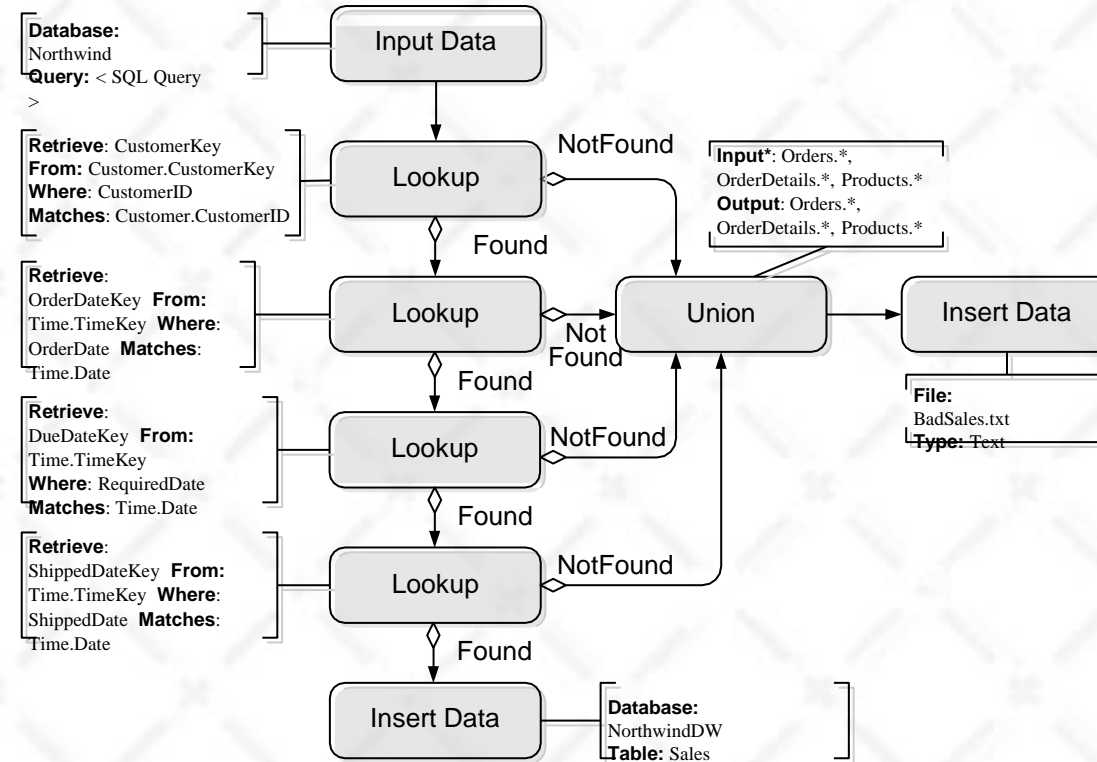
- ◆ The input table **Customers** is read from the operational database using an input data task
- ◆ **Region** (optional) in **Customers** is actually a state name or a state code → the first exclusive gateway checks whether this attribute is null or not
 - If **Region** is not null, add new column **State** initialized with the values of **Region**
 - Otherwise, check if the (City, Country) pair matches a pair in **TempCities**, and retrieve the **State** attribute, creating a new column
- ◆ A second exclusive gateway over the new **State** column accounts for countries without states
- ◆ Then perform a union over the two flows
- ◆ Finally, perform the union of all flows, and add the column **CustomerKey** for the surrogate key initialized to null

Load of the Territories Bridge Table

- ◆ The input is an SQL query joining **EmployeeTerritories** and **Territories**
- ◆ Then, an update column task removes ('trims') the leading spaces from attribute **TerritoryDescription**
- ◆ The city key is then obtained with a lookup over **City** in the D
- ◆ Finally, **Territories** is populated with an insert data task



Load of the Sales Fact Table



Load of the Sales Fact Table

- ◆ Task performed once all the other ones done
- ◆ Columns for order line number, sales amount, and freight must be created (Add Column data tasks)
- ◆ The process starts with an input data task that obtains data from the operational database via the query:

```
SELECT O.CustomerID, EmployeeID AS EmployeeKey, O.OrderDate,  
       O.RequiredDate AS DueDate, O.ShippedDate,  
       ShipVia AS ShipperKey, P.ProductID AS ProductKey, P.SupplierID AS  
       SupplierKey, O.OrderID AS OrderNo, ROW NUMBER() OVER  
       (PARTITION BY D.OrderID  
        ORDER BY D.ProductID) AS OrderLineNo, D.UnitPrice, Quantity, Discount, D.UnitPrice  
       * (1-Discount) * Quantity AS SalesAmount, O.Freight/COUNT(*) OVER (PARTITION BY  
       D.OrderID) AS Freight  
FROM      Orders O, OrderDetails D, Products P  
WHERE O.OrderID = D.OrderID AND D.ProductID = P.ProductID
```

- ◆ A sequence of lookups follows, which obtains the missing foreign keys for the dimension tables
- ◆ Finally, the fact table is loaded with the data retrieved

Chapter 6: Extraction, Transformation, and Loading

Outline

- ◆ Business Process Modeling Notation
- ◆ Conceptual ETL Design using BPMN
- ◆ Conceptual Design of the Northwind ETL

Integration Services and Kettle

- ◆ The Northwind ETL process in Integration Services
- ◆ The Northwind ETL process in Kettle

Integration Services

- ◆ SQL Server component to perform data migration tasks, and implement and execute ETL processes
- ◆ **Components of Integration Services**
 - **Package**: A workflow containing a collection of tasks executed in an orderly fashion
 - A package consists of a control flow and, optionally, one or more data flows
 - **Control flow**: three kinds of elements
 - * **Tasks**: Individual units of work that provide functionality to a package
 - Tasks: data flow tasks, data preparation tasks, Analysis Services tasks, workflow tasks
 - * **Containers**: Group tasks logically into units of work, and are used to define variables and events
 - Ex: **Sequence Container** and **For Loop Container**
 - * **Precedence constraints**: Connect tasks, containers, and executables defining execution order
- ◆ Creating a control flow in Integration Services requires:
 - Adding containers
 - Adding tasks
 - Connecting containers and tasks, using precedence constraints
 - Adding connection managers, when a task connects to a data source

Integration Services: Data Flows

- ◆ Extract data into memory, transform them, and write them to a destination
- ◆ Three kinds of components:
 - **Sources**: Extract data from data stores (OLE DB data sources, Excel files, flat files, and XML files, among other)
 - **Transformations**: Modify, summarize, and clean data (split, divert, or merge the flow)
 - * Example: **Conditional Split**, **Copy Column**, and **Aggregate**.
 - **Destinations**: Load data into data stores or create in-memory datasets
- ◆ Creating a data flow includes the following steps
 - Adding one or more sources
 - Adding the transformations to satisfy the package requirements
 - Connecting data flow components
 - Adding one or more destinations to load data into data stores
 - Configuring error outputs
 - Including annotations to document the data flow

◆ Main components:

- **Transformations**: Logical tasks consisting in **steps** connected by **hops**, essentially data flows to extract, transform, and load data
 - * **Steps**: Perform a specific tasks, e.g., reading data from a file, filtering rows, writing to a database
 - Steps grouped according to their function, such as input, output, scripting, etc.
 - * **Hops**: Data paths connecting steps to each other, so records can pass from one step to another
 - **Jobs**: Workflows that orchestrate the individual pieces of functionality implementing an entire ETL process
 - Jobs are composed of:
 - * **Jobs entries**: Primary building blocks of a job, correspond to the steps in data transformations
 - * **Jobs hops**: Specify the execution order of job entries and the conditions
 - * **Jobs settings**: Options that control the behavior of a job and the logging method
- ◆ **Important**: loops are not allowed in transformations, but allowed in jobs

◆ Kettle is composed of the following components:

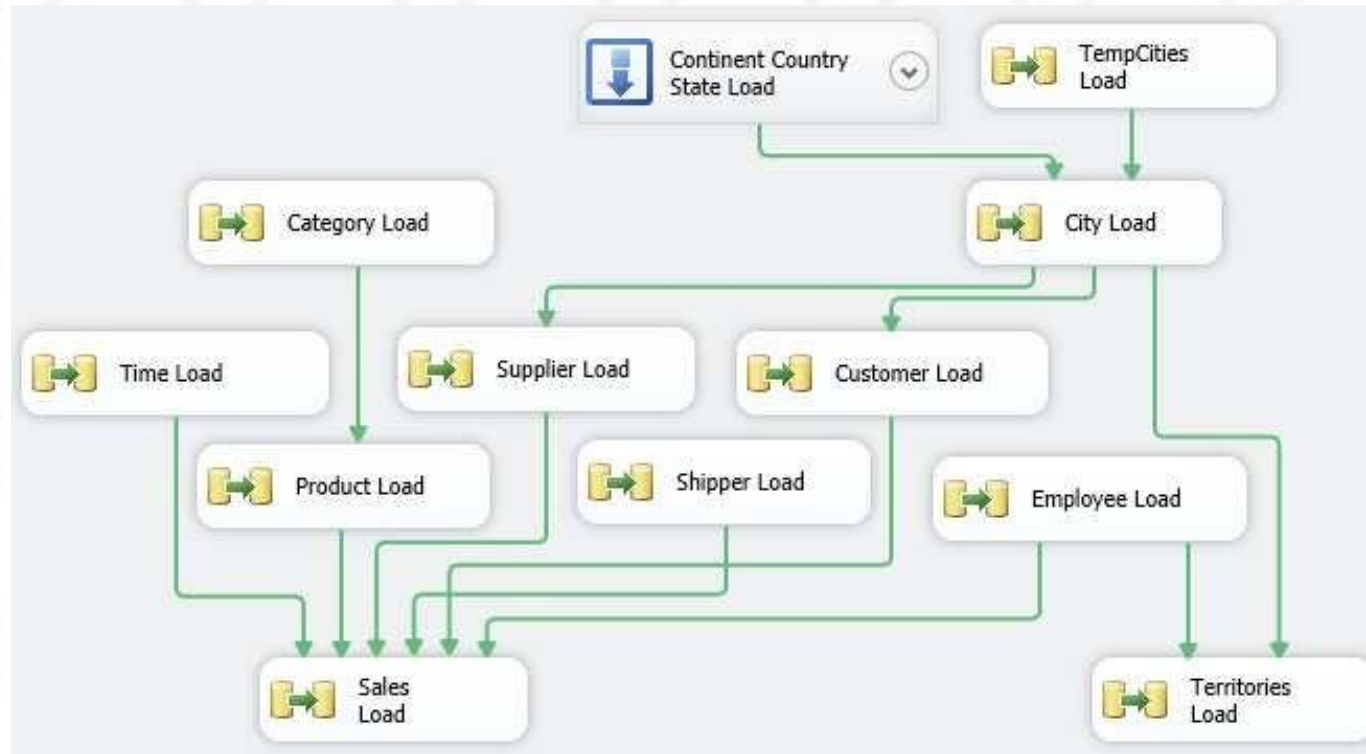
- **Data Integration Server**: Performs the actual data integration tasks
 - * Executes jobs and transformations
 - * Defines and manages security
 - * Provides content management
 - * Schedules and monitor activities
- **Spoon**: A graphical user interface for designing jobs and transformations
 - * Transformations can be executed locally within Spoon, or in the Data Integration Server
- **Pan**: A standalone command line tool for executing transformations
- **Kitchen**: A standalone command line tool for executing jobs Jobs are usually scheduled to run in batch mode at regular intervals.
- **Carte**: A lightweight server for running jobs and transformations on a remote host

Chapter 6: Extraction, Transformation, and Loading

Outline


- ◆ Business Process Modeling Notation
- ◆ Conceptual ETL Design using BPMN
- ◆ Conceptual Design of the Northwind ETL
- ◆ Integration Services and Kettle
- ◆ The Northwind ETL process in Integration Services
- ◆ The Northwind ETL process in Kettle

The Northwind ETL in Integration Services




Data Flow Tasks

- ◆ Many data flow tasks are simple
- ◆ These data flow tasks are composed of an OLE DB Source task that reads the table from the operational database and an OLE DB Destination task that receives the output and stores it in the DW
- ◆ Loading the **Category** dimension table

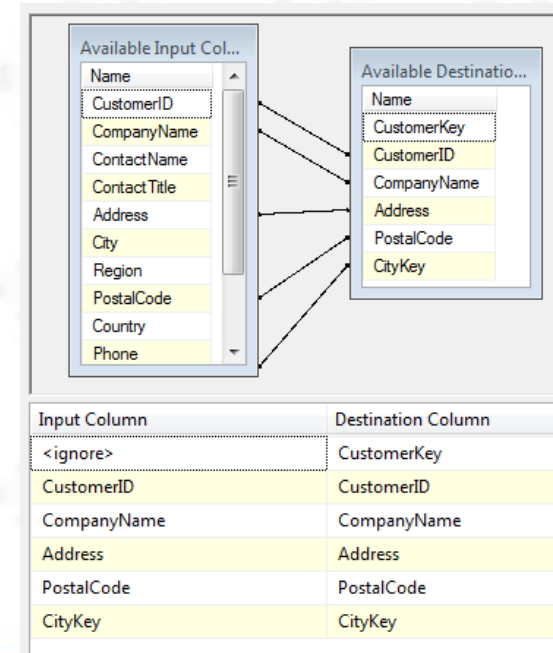
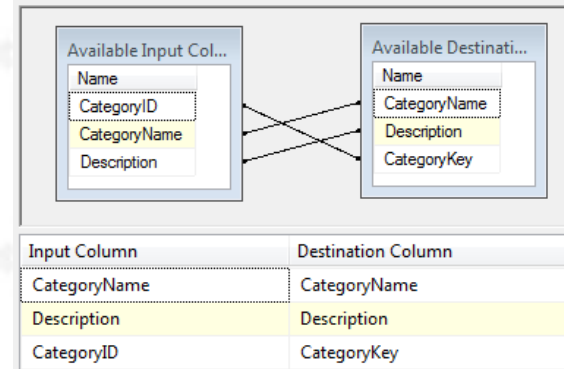


- ◆ Similar data flows are used for loading the **Product**, **Shipper**, and **Employee** tables
- ◆ Also straightforward is the data flow that loads the **Time** dimension from the source Excel file after a data conversion




Keys in the Data Warehouse

- ◆ Keys of the operational database are reused in the DW where dimensions do not have an alternate key
- ◆ For example, for table **Category** we reuse **CategoryID** as the key in the DW (**CategoryKey**)
- ◆ For table **Customer** the **CustomerID** key is an **CustomerAltKey** column in the DW
- ◆ A new value for **CustomerKey** is generated during the insert in the DW
- ◆ Mappings of the source and destination columns depending on the reuse of the key



Load of the Continent → Country → State Hierarchy

- ◆ Sequence container used for the three data flows that load the tables of the hierarchy



- ◆ Load of the **Continent** level

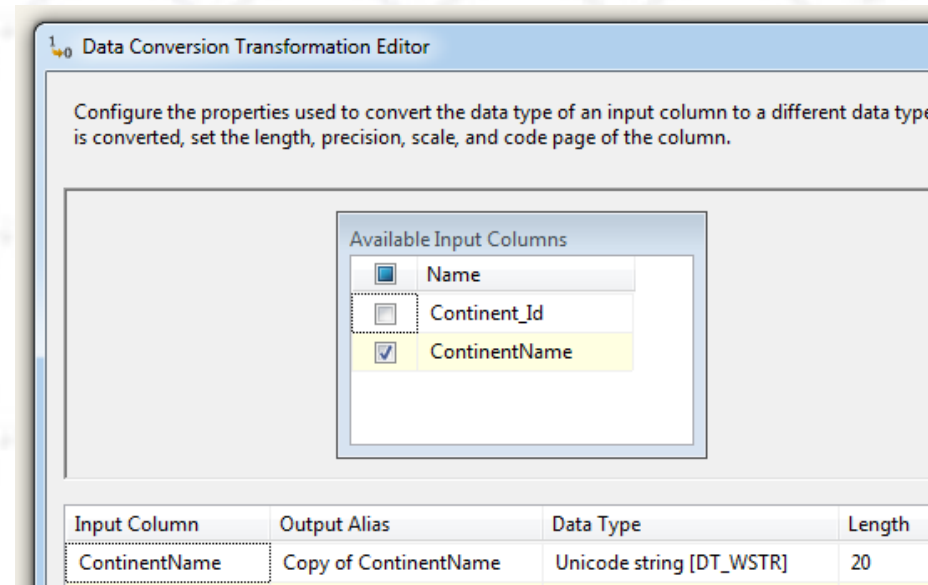


- ◆ Load of the **Country** level



- ◆ First produce a key to reference **Continent** from **Country**
- ◆ Data conversion tasks detailed in the next slide
- ◆ In the data flow that loads **Country** a merge join obtains the **ContinentName** for a given **Country**

Conversion of the Data Input from the XML File



- ◆ A data conversion transforms the data types from the XML file into the data types of the database
- ◆ The **ContinentName** read from the XML file is by default of length 255, and it is converted into a string of length 20

Load of the TempCities Table

- ◆ **TempCities**: A temporary table needed to load the geographic hierarchy associated to dimensions **Customer** and **Supplier**

- TempCities: Obtained from the text file **Cities.txt**

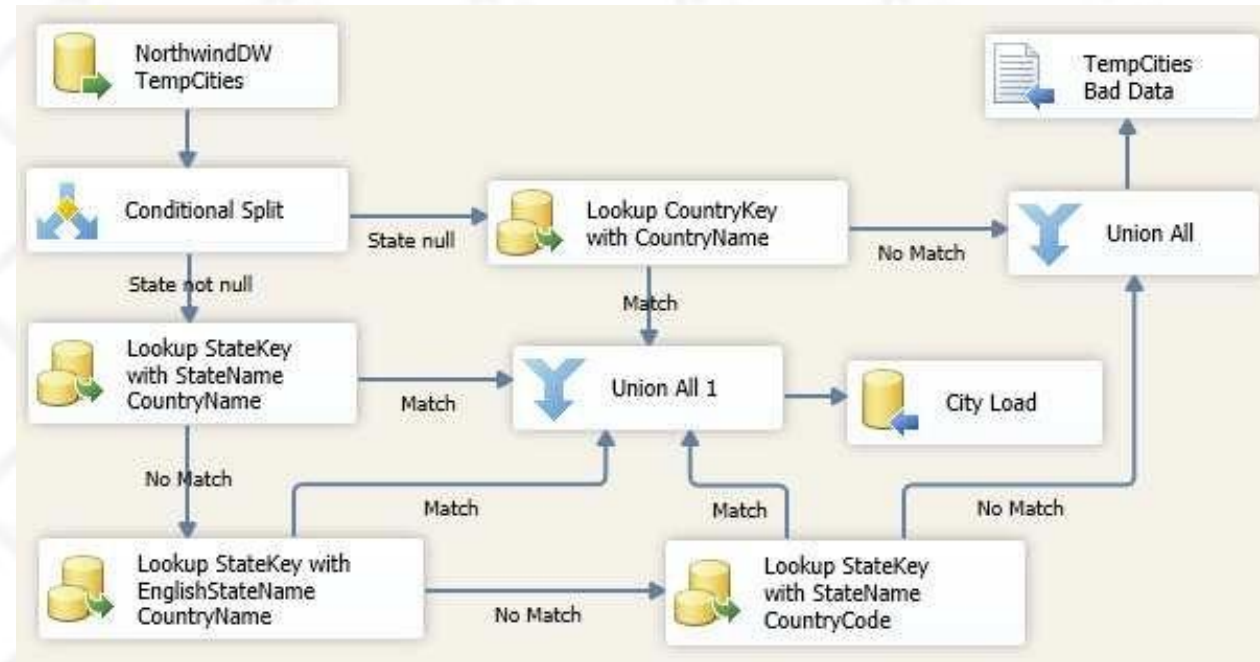
- ◆ Structure of the temporary table

| TempCities |
|------------|
| City |
| State |
| Country |

- ◆ We assume that this table already exists in the database
- ◆ A data conversion transformation is needed to transform the default types obtained from the text file into the database types

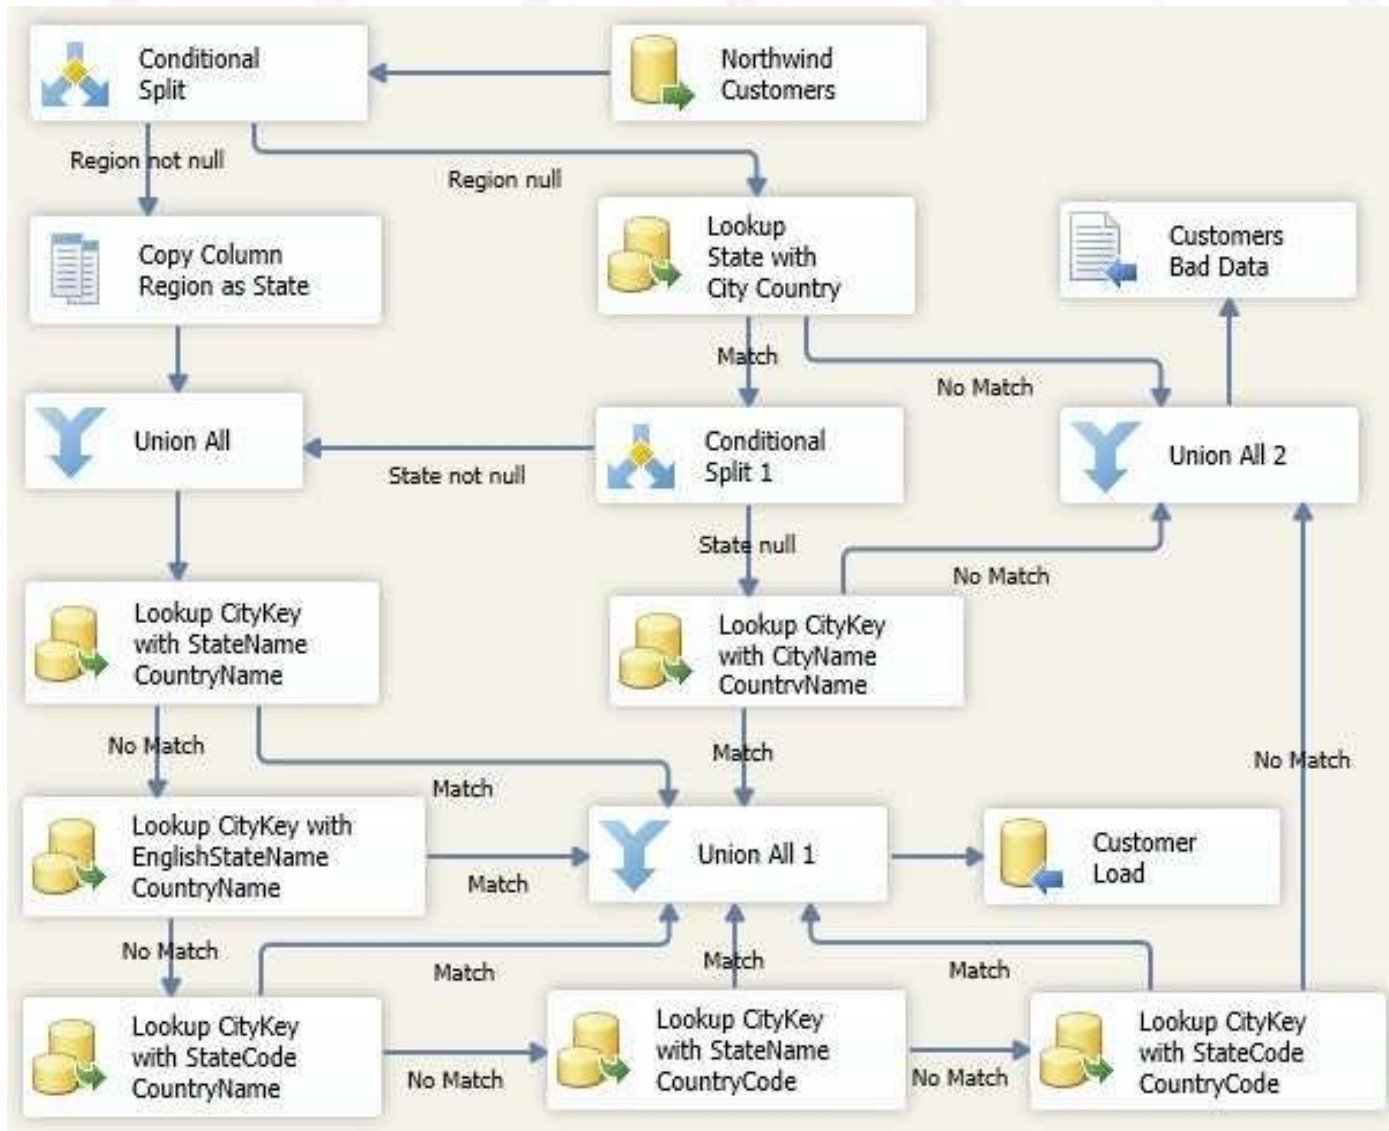


Load of the City Level



- ◆ The data flow associates to each city in **TempCities**, either a **StateKey** or a **CountryKey**, depending on whether or not the corresponding country is divided in states. For this:
 - The conditional split tests if the **State** is null or not
 - If so, a lookup is needed for obtaining the **CountryKey**

Load of the Customer Level



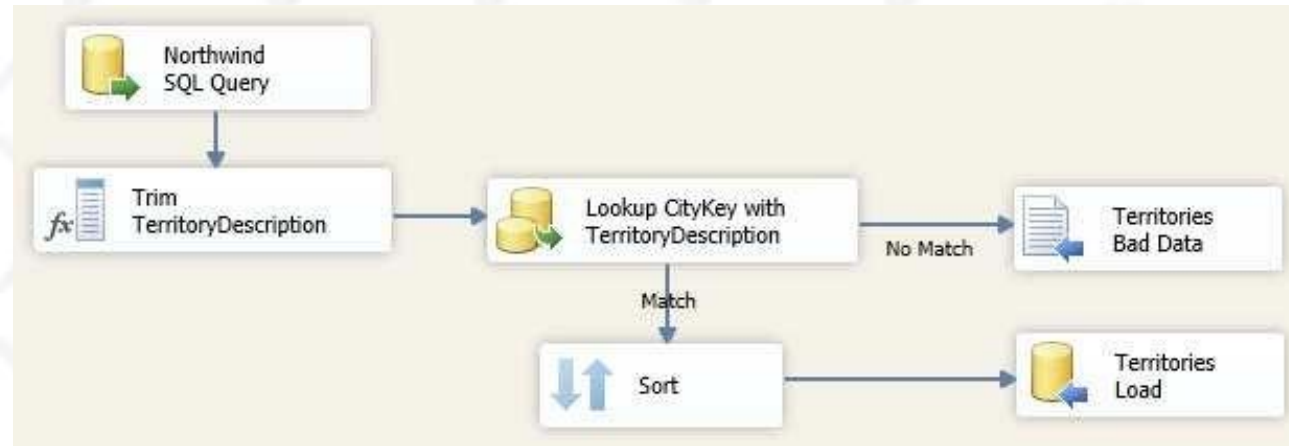
Load of the Customer Level

- ◆ Starts with a conditional split: if a customer has a null value in **Region**, a lookup adds a column **State**
by matching **City** and **Country** from **Customers** with **City** and **Country** from **TempCities**
- ◆ The value **State** obtained may be null for countries without states ⇒ a conditional split is needed
- ◆ If state is null, then a lookup tries to find a **CityKey** matching values of **City** and **Country** in a lookup table built as a join between **City** and **Country**

```
SELECT CityKey, CityName, CountryName  
FROM City C JOIN Country T ON C.CountryKey = T.CountryKey
```
- ◆ For customers with nonnull **Region**, the values of this column are copied into a new column **State**
- ◆ Then, 5 lookup tasks are needed, where each one tries to match a couple of values of **State** and **Country** to values in the lookup table built as a join between the **City**, **State**, and **Country** tables:

```
SELECT C.CityKey, C.CityName, S.StateName,  
       S.EnglishStateName, S.StateCode, T.CountryName,  
       T.CountryCode  
FROM City C JOIN State S ON C.StateKey = S.StateKey JOIN Country  
       T ON S.CountryKey = T.CountryKey
```

Load of the Territories Fact Table



- ◆ The data flow task starts with an OLE DB Source task, an SQL query:
SELECT E.*, TerritoryDescription
FROM EmployeeTerritories E JOIN Territories T ON E.TerritoryID = T.TerritoryID
- ◆ Continues with a derived column transformation that removes the trailing spaces in the values of **TerritoryDescription**
- ◆ A lookup transformation searches the corresponding values of **CityKey** in **City**
- ◆ Then a sort transformation removes duplicates

Load of the Sales Fact Table



- ◆ The first OLE DB Source task is a query that combines data from the operational DB and the DW (see the query in next slide)
- ◆ Then, a conditional split transformation task selects the records obtained in the query containing a null value in the columns **CustomerKey** or **ShippedDateKey**, and stores them in a flat file.
- ◆ The correct records are inserted in the data warehouse.

Load of the Sales Fact Table

- ◆ Query of the OLE DB source task

SELECT

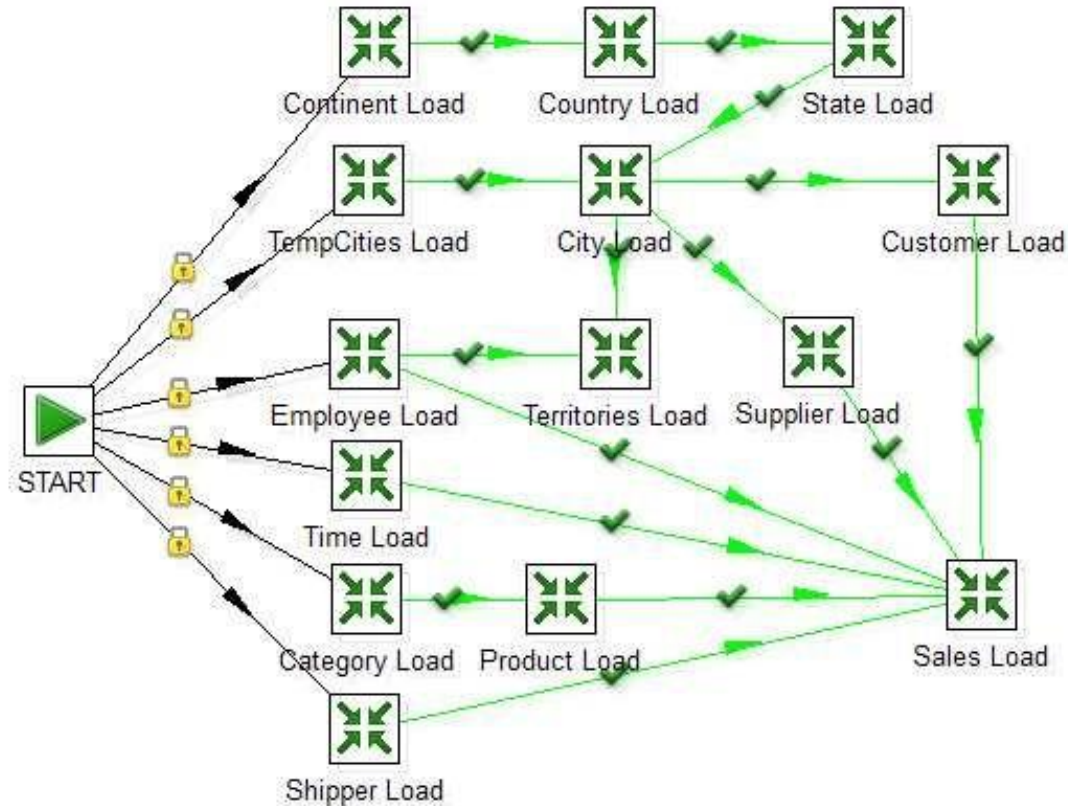
```
( SELECT CustomerKey FROM dbo.Customer C
  WHERE C.CustomerID = O.CustomerID) AS CustomerKey, EmployeeID AS
EmployeeKey,
( SELECT TimeKey FROM dbo.Time T
  WHERE T.Date = O.OrderDate) AS OrderDateKey,
( SELECT TimeKey FROM dbo.Time T
  WHERE T.Date = O.RequiredDate) AS DueDateKey,
( SELECT TimeKey FROM dbo.Time T
  WHERE T.Date = O.ShippedDate) AS ShippedDateKey, ShipVia AS
ShipperKey, P.ProductID AS ProductKey, SupplierID AS SupplierKey,
O.OrderID AS OrderNo,
CONVERT(INT, ROW NUMBER() OVER (PARTITION BY D.OrderID
  ORDER BY D.ProductID)) AS OrderLineNo, D.UnitPrice, Quantity, Discount, CONVERT(MONEY,
D.UnitPrice * (1-Discount) * Quantity) AS SalesAmount, CONVERT(MONEY, O.Freight/COUNT(*)
OVER (PARTITION BY D.OrderID)) AS Freight
FROM Northwind.dbo.Orders O, Northwind.dbo.OrderDetails D, Northwind.dbo.Products P
WHERE O.OrderID = D.OrderID AND D.ProductID = P.ProductID
```

Chapter 6: Extraction, Transformation, and Loading

Outline


- ◆ Business Process Modeling Notation
- ◆ Conceptual ETL Design using BPMN
- ◆ Conceptual Design of the Northwind ETL
- ◆ Integration Services and Kettle
- ◆ The Northwind ETL process in Integration Services
- ◆ The Northwind ETL process in Kettle

Overall View of the ETL Process




Load of the Category and Time Dimension Tables


- ◆ Loading the **Category** dimension table is similar to the data flow in Integration Services



- ◆ Loading the **Time** dimension table: Specified in the transformation step that reads the CSV file




Load of the Employee Dimension Table




- ◆ Transformation requires a **closure table** containing the transitive closure of **Supervision** hierarchy
- ◆ After reading the **Employees** table the rows read are sent in parallel to the steps that load the **Employee** and the **EmployeeClosure** tables

Load of the Continent and Country Levels

- ◆ Load of the **Continent** level

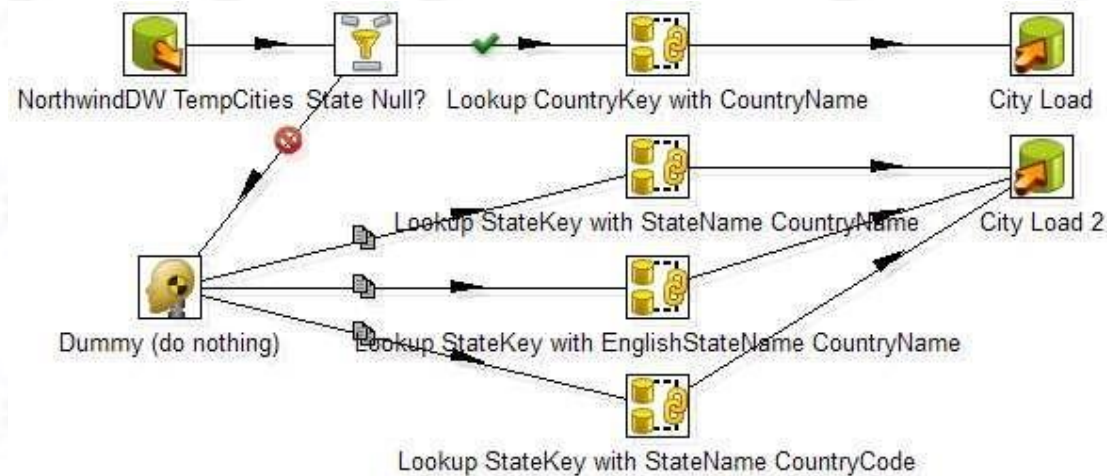


- ◆ With respect to IS, the conversion task is not required in Kettle
- ◆ Load of the **Country** level



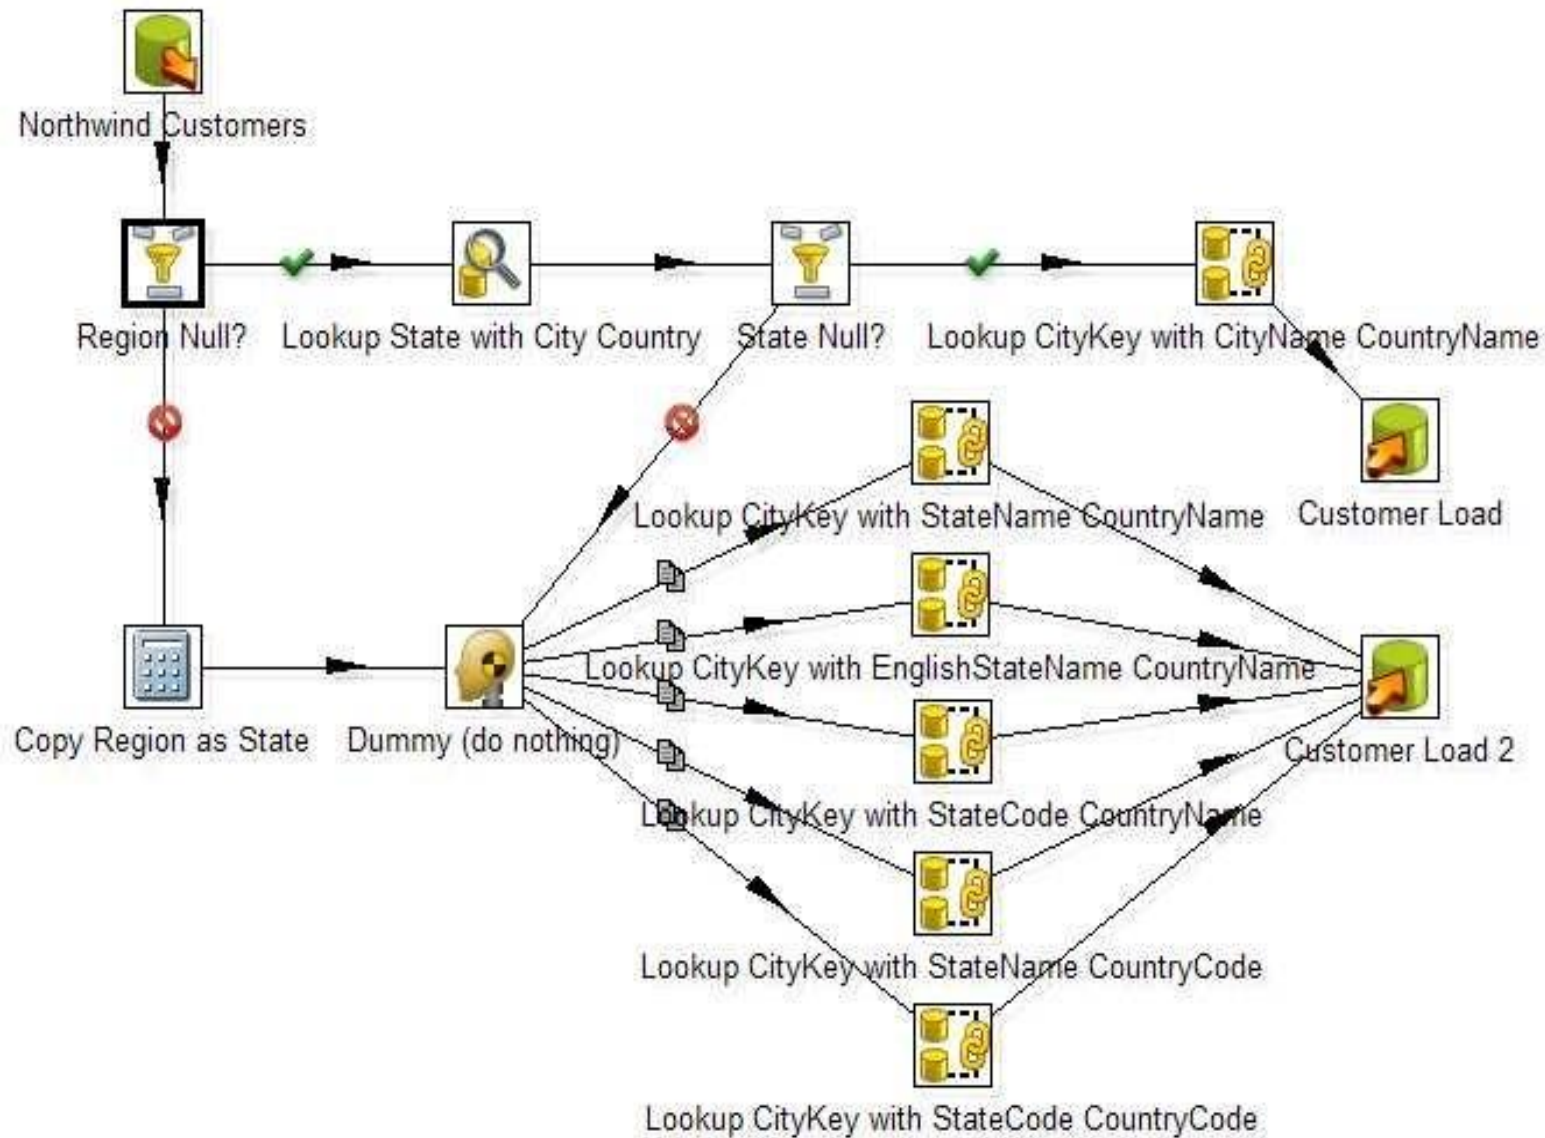
- ◆ With respect to IS, in Kettle we can find the **ContinentName** associated to a **Country** using an XPath expression, as in the conceptual design

Load of the City Level



- ◆ Significant *differences* with IS
- ◆ No possible to cascade lookup steps in Kettle as it is done with lookup tasks in IS
- ◆ Cascade lookups must be implemented as a collection of parallel flows.
- ◆ In Kettle we do not have tasks that load records for which a lookup value was not found in a text file
- ◆ The rows that do not have a null value sent in parallel to all the subsequent lookup tasks
- ◆ A dummy task is needed
- ◆ In Kettle there is no need to explicitly include a union task but all fields in the input flows have the same name → need one step for **CountryKey**, and other for **StateKey**

Load of the Customer Level




Load of the Customer Level

- ◆ Two different steps for performing lookups (different icons):
 - One that looks for **State**, and the other ones that look for **CityKey**
 - The former lookup type looks for values *in a single table* and sends all rows to the output flow
 - The second type of lookup looks for values *in an SQL query* and only sends to the output stream the rows with matching value
- ◆ A dummy task is needed in Kettle
- ◆ Dummy step sends the input rows to all subsequent lookup tasks
- ◆ SQL query used in the lookup step that looks for **CityKey** with **StateName** and **CountryName**

```
SELECT C.CityKey
FROM City C JOIN State S ON C.StateKey = S.Statekey JOIN Country T
ON S.CountryKey = T.CountryKey
WHERE ? = CityName AND ? = StateName AND ? = CountryName
```

Load of the Territories Fact Table



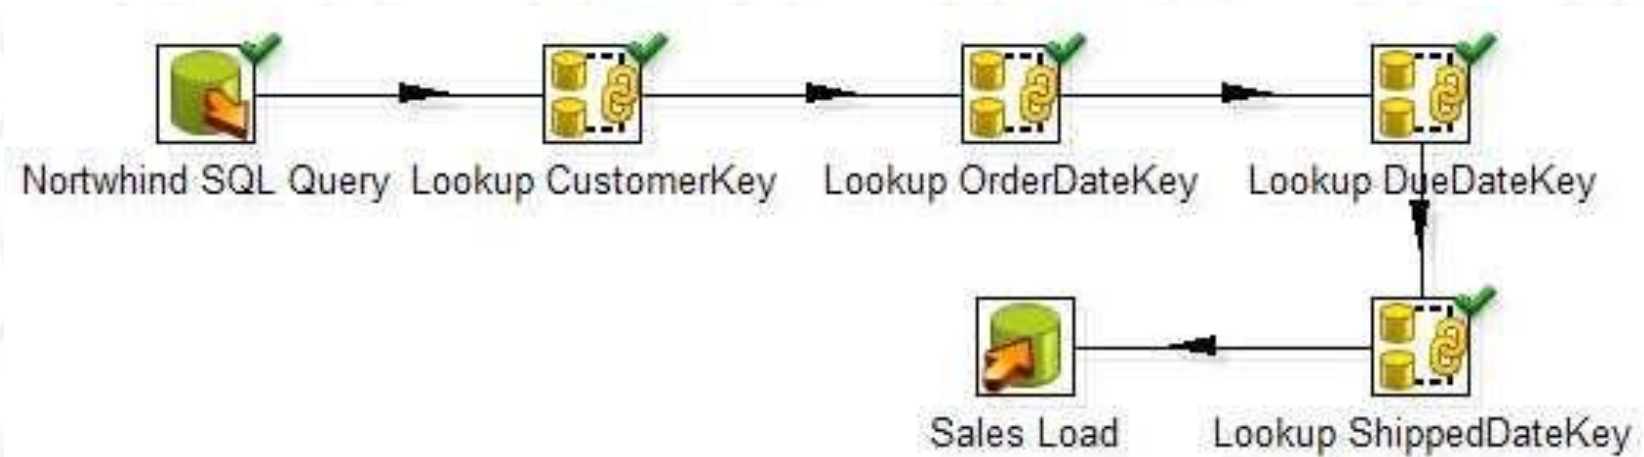
- ◆ Flow starts by obtaining the assignment of employees to territories from the Northwind database using an SQL query
- ◆ In Kettle there is no step that removes the trailing spaces in the **TerritoryDescription** column
- ◆ This was taken into account in the SQL query of the subsequent lookup step:
SELECT CityKey FROM City
WHERE TRIM(?) = CityName
- ◆ After the lookup of the **CityKey**, Kettle requires a sort but this does not remove duplicates, like in IS

Load of the Territories Fact Table

- ◆ The flow starts by obtaining values from the following SQL query addressed to the Northwind database, like in the conceptual design

```
SELECT O.CustomerID, EmployeeID AS EmployeeKey,  
       O.OrderDate, O.RequiredDate, O.ShippedDate, ShipVia  
       AS ShipperKey, P.ProductID AS ProductKey,  
       P.SupplierID AS SupplierKey, O.OrderID AS OrderNo,  
       ROW NUMBER() OVER (PARTITION BY D.OrderID  
                           ORDER BY D.ProductID) AS OrderLineNo,  
       D.UnitPrice, Quantity, Discount,  
       D.UnitPrice * (1-Discount) * Quantity AS SalesAmount,  
       O.Freight/COUNT(*) OVER (PARTITION BY D.OrderID) AS  
       Freight  
FROM    Orders O, OrderDetails D, Products P  
WHERE   O.OrderID = D.OrderID AND D.ProductID = P.ProductID
```


Load of the Sales Fact Table



- ◆ In IS it is possible to query both the Northwind operational database and the Northwind data warehouse in a single query
- ◆ Not possible in PostgreSQL
- ◆ Thus, additional lookup steps are needed in Kettle for obtaining the surrogate keys
- ◆ Additional task needed in IS for removing the records with null values for surrogate keys
- ◆ These are automatically removed in the lookup steps in Kettle

Enjoy the Course...!